

Kubernetes 1.33完全実装ガイド

Proxmox + Terraform + Ansible + helmfile統合構成

対象環境: 開発環境（本番環境への拡張可能）

構成: Control Plane 2台 + Worker 3台

最終更新: 2025年10月

目次

- 1. [アーキテクチャ概要](#)
- 2. [バージョン一覧](#)
- 3. [Proxmox + Terraform構成](#)
- 4. [Ansible自動構築](#)
- 5. [Cilium CNI設定（kube-proxy完全置換）](#)
- 6. [helmfile環境管理](#)
- 7. [全コンポーネント設定詳細](#)
- 8. [メトリクス可視化設定](#)
- 9. [デプロイ手順](#)
- 10. [トラブルシューティング](#)

1. アーキテクチャ概要

1.1 全体構成図



Proxmox VE

Terraform管理VM群

CP-1	CP-2	Control Plane
4vCPU	4vCPU	
8GB RAM	8GB RAM	

Worker-1	Worker-2	Worker-3
8vCPU	8vCPU	8vCPU
16GB RAM	16GB RAM	16GB RAM



Ansible自動構築（kubeadm）



Kubernetes 1.33クラスター

- Cilium CNI（kube-proxy完全置換）
- LB IPAM（MetalLB代替）
- Hubble可視化
- Gateway API（HTTP/HTTPS/TLS/gRPC）

- Traefik Ingress
- TCP/UDPルーティング
 - 全UIへのアクセス提供

監視スタック

- Prometheus（メトリクス収集）		
- Grafana（可視化）		
- Vector（ログ・メトリクス集約）		
- Jaeger（トレーシング）		
ストレージ・管理		
- MinIO（S3互換）		
- ArgoCD（GitOps）		
- Kubernetes Dashboard		

1.2 ネットワーク設計

用途	CIDR	例
Proxmox管理ネットワーク	192.168.1.0/24	-
Control Plane	192.168.1.10-11	CP-1: .10, CP-2: .11
Worker	192.168.1.20-22	W-1: .20, W-2: .21, W-3: .22
API VIP	192.168.1.100	HAProxy/keepalived
Cilium LB IPAM Pool	192.168.100.0/24	LoadBalancer Service用
Pod Network	10.244.0.0/16	Cilium管理
Service Network	10.96.0.0/12	ClusterIP範囲

1.3 重要な設計判断

✅ 採用する技術

- **Cilium Gateway API:** HTTP/HTTPS/gRPC/TLS専用（本番実績あり）
- **Traefik Ingress:** TCP/UDP用（Gateway APIのTCPRoute/UDPRouteは代替）
- **kube-proxy完全置換:** eBPFベースの高性能化
- **LB IPAM:** MetalLB不要、Cilium統合機能を使用
- **helmfile:** 環境切り替え + GitOps準備

❌ 使用しない技術

- **Cilium Gateway API for TCP/UDP:** 2025年10月時点で未実装（GitHub Issue #21929オープン中）
- **MetalLB:** Cilium LB IPAMで代替
- **Istio/Linkerd:** 開発環境では過剰

2. バージョン一覧

2.1 推奨バージョン

コンポーネント	バージョン	リリース日	サポート期限	Helmチャート
Kubernetes	1.33.5	2025年1月	2026年6月	-
Cilium	1.18.2	2025年9月	LTS	cilium/cilium:1.18.2
Traefik	3.5.0	2025年10月	-	traefik/traefik:37.1.2
kube-prometheus-stack	78.3.0	2025年10月	-	prometheus-community/kube-prometheus-stack:78.3.0
Grafana	11.x	2025年	-	(stackに含む)
Jaeger Operator	1.65.0	2025年	-	jaegertracing/jaeger-operator:2.57.0
Vector	0.46.0	2025年	-	vector/vector:0.46.0
MinIO Operator	7.1.1	2025年	-	minio-operator/operator:7.1.1
ArgoCD	2.13.x	2025年	-	argo/argo-cd:8.6.2
Kubernetes Dashboard	7.13.0	2025年	-	kubernetes-dashboard/kubernetes-dashboard:7.13.0

2.2 互換性マトリックス

コンポーネント	K8s 1.31	K8s 1.32	K8s 1.33	K8s 1.34
Cilium 1.18	✓	✓	✓	✓
Traefik 3.5	✓	✓	✓	✓
kube-prometheus-stack 78.x	✓	✓	✓	✓
ArgoCD 2.13	✓	✓	✓	✓

3. Proxmox + Terraform構成

3.1 ディレクトリ構造



```
infrastructure/
├── terraform/
│   ├── main.tf
│   ├── variables.tf
│   ├── outputs.tf
│   ├── terraform.tfvars
│   └── modules/
│       ├── control-plane/
│       │   ├── main.tf
│       │   └── variables.tf
│       └── worker/
│           ├── main.tf
│           └── variables.tf
├── ansible/
│   ├── inventory/
│   │   └── hosts.ini
│   ├── group_vars/
│   │   ├── all.yml
│   │   └── k8s-cluster.yml
│   ├── playbooks/
│   │   ├── 01-prerequisites.yml
│   │   ├── 02-container-runtime.yml
│   │   ├── 03-kubernetes-install.yml
│   │   └── 04-cluster-init.yml
│   └── templates/
│       └── kubeadm-config.yaml.j2
└── helmfile/
    ├── helmfile.yaml
    ├── environments/
    │   ├── dev.yaml
    │   ├── staging.yaml
    │   └── prod.yaml
    └── values/
        ├── dev/
        ├── staging/
        └── prod/
```

3.2 Terraform設定

terraform/main.tf:



hcl

```
terraform {  
  required_version = ">= 1.9"  
  required_providers {  
    proxmox = {  
      source = "Telmate/proxmox"  
      version = "~> 3.0"  
    }  
  }  
}
```

```
provider "proxmox" {  
  pm_api_url    = var.proxmox_api_url  
  pm_user       = var.proxmox_user  
  pm_password   = var.proxmox_password  
  pm_tls_insecure = true  
}
```

Control Plane VMs

```
module "control_plane" {  
  source = "./modules/control-plane"  
  
  count      = 2  
  vm_name    = "k8s-cp-${count.index + 1}"  
  target_node = var.proxmox_node  
  vm_id      = 200 + count.index  
  
  cores      = 4  
  sockets    = 1  
  memory     = 8192  
  disk_size  = "100G"  
  
  network_bridge = "vmbr0"  
  ip_address     = "192.168.1.${10 + count.index}"  
  gateway        = var.gateway  
  
  ssh_keys    = var.ssh_public_keys  
  cloud_init_image = var.cloud_init_image  
}
```

Worker VMs

```
module "worker" {  
  source = "./modules/worker"  
  
  count = 3
```

```
vm_name      = "k8s-worker-${count.index + 1}"
target_node  = var.proxmox_node
vm_id        = 210 + count.index

cores        = 8
sockets      = 1
memory       = 16384
disk_size    = "200G"

network_bridge = "vmbr0"
ip_address     = "192.168.1.${20 + count.index}"
gateway       = var.gateway

ssh_keys      = var.ssh_public_keys
cloud_init_image = var.cloud_init_image
}
```

Output

```
output "control_plane_ips" {
  value = [for vm in module.control_plane : vm.ip_address]
}

output "worker_ips" {
  value = [for vm in module.worker : vm.ip_address]
}
```

terraform/variables.tf:



hcl

```
variable "proxmox_api_url" {  
  description = "Proxmox API URL"  
  type        = string  
  default     = "https://proxmox.example.com:8006/api2/json"  
}
```

```
variable "proxmox_user" {  
  description = "Proxmox user"  
  type        = string  
}
```

```
variable "proxmox_password" {  
  description = "Proxmox password"  
  type        = string  
  sensitive   = true  
}
```

```
variable "proxmox_node" {  
  description = "Proxmox node name"  
  type        = string  
  default     = "pve"  
}
```

```
variable "gateway" {  
  description = "Network gateway"  
  type        = string  
  default     = "192.168.1.1"  
}
```

```
variable "ssh_public_keys" {  
  description = "SSH public keys"  
  type        = list(string)  
}
```

```
variable "cloud_init_image" {  
  description = "Cloud-init template name"  
  type        = string  
  default     = "ubuntu-22.04-cloud"  
}
```

terraform/terraform.tfvars:



hcl

```
proxmox_api_url = "https://192.168.1.5:8006/api2/json"
proxmox_user    = "root@pam"
proxmox_node    = "pve"
gateway        = "192.168.1.1"

ssh_public_keys = [
  "ssh-rsa AAAAB3NzaC1yc2EA... user@laptop"
]

cloud_init_image = "ubuntu-22.04-cloud-init"
```

terraform/modules/control-plane/main.tf:



hcl

```
resource "proxmox_vm_qemu" "control_plane" {
  name      = var.vm_name
  target_node = var.target_node
  vmid      = var.vm_id

  clone    = var.cloud_init_image
  full_clone = true

  cores    = var.cores
  sockets = var.sockets
  memory   = var.memory

  disk {
    size    = var.disk_size
    type    = "scsi"
    storage = "local-lvm"
    iothread = 1
  }

  network {
    model = "virtio"
    bridge = var.network_bridge
  }

  # Cloud-init
  os_type = "cloud-init"
  ipconfig0 = "ip=${var.ip_address}/24,gw=${var.gateway}"

  sshkeys = join("\n", var.ssh_keys)

  lifecycle {
    ignore_changes = [
      network,
    ]
  }

  tags = "kubernetes,control-plane"
}

output "ip_address" {
  value = var.ip_address
}
```

3.3 Terraform実行



bash

```
# 初期化
cd terraform/
terraform init

# 計画確認
terraform plan

# 実行
terraform apply -auto-approve

# 出力確認
terraform output
```

4. Ansible自動構築

4.1 インベントリ

ansible/inventory/hosts.ini:



ini

```
[control_plane]
```

```
k8s-cp-1 ansible_host=192.168.1.10 ansible_user=ubuntu
```

```
k8s-cp-2 ansible_host=192.168.1.11 ansible_user=ubuntu
```

```
[workers]
```

```
k8s-worker-1 ansible_host=192.168.1.20 ansible_user=ubuntu
```

```
k8s-worker-2 ansible_host=192.168.1.21 ansible_user=ubuntu
```

```
k8s-worker-3 ansible_host=192.168.1.22 ansible_user=ubuntu
```

```
[k8s_cluster:children]
```

```
control_plane
```

```
workers
```

```
[all:vars]
```

```
ansible_python_interpreter=/usr/bin/python3
```

```
ansible_ssh_common_args='-o StrictHostKeyChecking=no'
```

4.2 Group Variables

ansible/group_vars/all.yml:



yaml

Kubernetes設定

kubernetes_version: "1.33.5"

kubernetes_cni: "cilium"

skip_kube_proxy: true

Container Runtime

container_runtime: "containerd"

containerd_version: "1.7.x"

Network設定

pod_network_cidr: "10.244.0.0/16"

service_cidr: "10.96.0.0/12"

cluster_dns: "10.96.0.10"

Control Plane

control_plane_endpoint: "192.168.1.100:6443" # VIPまたはDNS

api_server_cert_sans:

- "k8s-api.example.com"

- "192.168.1.100"

- "192.168.1.10"

- "192.168.1.11"

Cilium設定

cilium_version: "1.18.2"

cilium_lb_ipam_pool: "192.168.100.0/24"

4.3 Playbook 1: 前提条件

ansible/playbooks/01-prerequisites.yml:



yaml

- **name:** Kubernetes前提条件セットアップ

hosts: k8s_cluster

become: yes

tasks:

- **name:** システムアップデート

apt:

update_cache: yes

upgrade: dist

when: ansible_os_family == "Debian"

- **name:** 必要パッケージインストール

apt:

name:

- apt-transport-https

- ca-certificates

- curl

- gnupg

- lsb-release

- socat

- conntrack

- ipset

state: present

- **name:** Swapの無効化

shell: |

swapoff -a

sed -i '/swap/d' /etc/fstab

- **name:** カーネルモジュールの永続化

copy:

dest: /etc/modules-load.d/k8s.conf

content: |

overlay

br_netfilter

- **name:** カーネルモジュールのロード

modprobe:

name: "{{ item }}"

state: present

loop:

- overlay

- br_netfilter

- **name:** sysctl/パラメータ設定
copy:
dest: /etc/sysctl.d/k8s.conf
content: |
net.bridge.bridge-nf-call-iptables = 1
net.bridge.bridge-nf-call-ip6tables = 1
net.ipv4.ip_forward = 1
- **name:** sysctl/パラメータ適用
command: sysctl --system
- **name:** ファイアウォール無効化（開発環境）
systemd:
name: ufw
state: stopped
enabled: no
ignore_errors: yes

4.4 Playbook 2: Container Runtime

ansible/playbooks/02-container-runtime.yml:



yaml

- **name:** containerdインストール

hosts: k8s_cluster

become: yes

tasks:

- **name:** Dockerリポジトリ追加

shell: |

install -m 0755 -d /etc/apt/keyrings

curl -fsSL https://download.docker.com/linux/ubuntu/gpg | gpg --dearmor -o /etc/apt/keyrings/doc

chmod a+r /etc/apt/keyrings/docker.gpg

echo "deb [arch=\$(dpkg --print-architecture) signed-by=/etc/apt/keyrings/docker.gpg] https://dowr

args:

creates: /etc/apt/sources.list.d/docker.list

- **name:** containerdインストール

apt:

name: containerd.io

state: present

update_cache: yes

- **name:** containerd設定ディレクトリ作成

file:

path: /etc/containerd

state: directory

- **name:** containerdデフォルト設定生成

shell: containerd config default > /etc/containerd/config.toml

args:

creates: /etc/containerd/config.toml

- **name:** SystemdCgroup有効化

lineinfile:

path: /etc/containerd/config.toml

regexp: 'SystemdCgroup = false'

line: 'SystemdCgroup = true'

state: present

- **name:** containerd再起動

systemd:

name: containerd

state: restarted

enabled: yes
daemon_reload: yes

4.5 Playbook 3: Kubernetesインストール

ansible/playbooks/03-kubernetes-install.yml:



yaml

- **name:** Kubernetesパッケージインストール
hosts: k8s_cluster
become: yes
tasks:
 - **name:** Kubernetesリポジトリキー追加
shell: |
mkdir -p /etc/apt/keyrings
curl -fsSL https://pkgs.k8s.io/core:/stable:/v1.33/deb/Release.key | gpg --dearmor -o /etc/apt/keyrin
args:
creates: /etc/apt/keyrings/kubernetes-apt-keyring.gpg
 - **name:** Kubernetesリポジトリ追加
copy:
dest: /etc/apt/sources.list.d/kubernetes.list
content: |
deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg] https://pkgs.k8s.io/core:/stable:/v1
 - **name:** Kubernetesパッケージインストール
apt:
name:
 - "kubernetes={{ kubernetes_version }}-*"
 - "kubeadm={{ kubernetes_version }}-*"
 - "kubectl={{ kubernetes_version }}-*"**state:** present
update_cache: yes
 - **name:** Kubernetesパッケージバージョン固定
dpkg_selections:
name: "{{ item }}"
selection: hold
loop:
 - kubelet
 - kubeadm
 - kubectl

4.6 Playbook 4: クラスタ初期化

ansible/playbooks/04-cluster-init.yml:



yaml

- name: 最初のControl Plane初期化
 - hosts: control_plane[0]
 - become: yes
 - vars:
 - kubeconfig_path: "{{ ansible_env.HOME }}/.kube/config"
 - tasks:
 - name: kubeadm設定ファイル生成
 - template:
 - src: ../templates/kubeadm-config.yaml.j2
 - dest: /root/kubeadm-config.yaml
 - name: クラスタ初期化チェック
 - stat:
 - path: /etc/kubernetes/admin.conf
 - register: k8s_init
 - name: クラスタ初期化
 - command: kubeadm init --config=/root/kubeadm-config.yaml --upload-certs
 - register: kubeadm_init_output
 - when: not k8s_init.stat.exists
 - name: 初期化ログ保存
 - copy:
 - content: "{{ kubeadm_init_output.stdout }}"
 - dest: /root/kubeadm-init.log
 - when: kubeadm_init_output.changed
 - name: kubeconfigディレクトリ作成
 - file:
 - path: "{{ ansible_env.HOME }}/.kube"
 - state: directory
 - mode: '0755'
 - name: kubeconfigコピー
 - copy:
 - src: /etc/kubernetes/admin.conf
 - dest: "{{ kubeconfig_path }}"
 - remote_src: yes
 - owner: "{{ ansible_user }}"
 - group: "{{ ansible_user }}"
 - mode: '0644'
 - name: Join Token取得

command: kubeadm token create

register: join_token

- **name:** CA証明書ハッシュ取得

shell: openssl x509 -pubkey -in /etc/kubernetes/pki/ca.crt | openssl rsa -pubin -outform der 2>/dev/n

register: ca_cert_hash

- **name:** Certificate Key取得

shell: kubeadm init phase upload-certs --upload-certs 2>/dev/null | tail -n 1

register: certificate_key

- **name:** Join情報を変数に保存

set_fact:

k8s_join_token: "{{ join_token.stdout }}"

k8s_ca_hash: "{{ ca_cert_hash.stdout }}"

k8s_cert_key: "{{ certificate_key.stdout }}"

- **name:** Helmインストール

shell: |

curl https://raw.githubusercontent.com/helm/helm/main/scripts/get-helm-3 | **bash**

args:

creates: /usr/local/bin/helm

- **name:** Ciliumインストール

shell: |

helm repo add cilium https://helm.cilium.io/ || **true**

helm repo update

**helm install cilium cilium/cilium **

--version {{ cilium_version }} \

**--namespace kube-system **

**--set kubeProxyReplacement=true **

**--set k8sServiceHost={{ control_plane_endpoint.split(':')[0] }} **

**--set k8sServicePort={{ control_plane_endpoint.split(':')[1] }} **

**--set hubble.enabled=true **

**--set hubble.relay.enabled=true **

**--set hubble.ui.enabled=true **

**--wait **

--timeout=10m

environment:

KUBECONFIG: "{{ kubeconfig_path }}"

- **name:** 2台目のControl Plane参加

hosts: control_plane[1:]

become: yes

serial: 1

tasks:

- name: Control Plane参加済みチェック

stat:

path: /etc/kubernetes/kubelet.conf

register: cp_joined

- name: Control Planeとして参加

command: >

kubeadm join {{ control_plane_endpoint }}

--token {{ hostvars[groups['control_plane'][0]]['k8s_join_token'] }}

--discovery-token-ca-cert-hash sha256:{{ hostvars[groups['control_plane'][0]]['k8s_ca_hash'] }}

--control-plane

--certificate-key {{ hostvars[groups['control_plane'][0]]['k8s_cert_key'] }}

--apiserver-advertise-address={{ ansible_default_ipv4.address }}

when: not cp_joined.stat.exists

- name: Worker参加

hosts: workers

become: yes

tasks:

- name: Worker参加済みチェック

stat:

path: /etc/kubernetes/kubelet.conf

register: worker_joined

- name: Workerとして参加

command: >

kubeadm join {{ control_plane_endpoint }}

--token {{ hostvars[groups['control_plane'][0]]['k8s_join_token'] }}

--discovery-token-ca-cert-hash sha256:{{ hostvars[groups['control_plane'][0]]['k8s_ca_hash'] }}

when: not worker_joined.stat.exists

- name: kubeconfig取得

hosts: control_plane[0]

become: yes

tasks:

- name: kubeconfigを取得

fetch:

src: /etc/kubernetes/admin.conf

dest: "{{ playbook_dir }}/../kubeconfig"

flat: yes

4.7 kubeadm設定テンプレート

ansible/templates/kubeadm-config.yaml.j2:



yaml

```
---
apiVersion: kubeadm.k8s.io/v1beta4
kind: InitConfiguration
nodeRegistration:
  name: "{{ ansible_hostname }}"
  criSocket: "unix:///var/run/containerd/containerd.sock"
  kubeletExtraArgs:
    - name: "node-ip"
      value: "{{ ansible_default_ipv4.address }}"
  taints:
    - key: "node-role.kubernetes.io/control-plane"
      effect: "NoSchedule"
localAPIEndpoint:
  advertiseAddress: "{{ ansible_default_ipv4.address }}"
  bindPort: 6443
{% if skip_kube_proxy %}
skipPhases:
  - addon/kube-proxy
{% endif %}
```

```
---
apiVersion: kubeadm.k8s.io/v1beta4
kind: ClusterConfiguration
kubernetesVersion: "v{{ kubernetes_version }}"
clusterName: "k8s-cluster"
controlPlaneEndpoint: "{{ control_plane_endpoint }}"
networking:
  serviceSubnet: "{{ service_cidr }}"
  podSubnet: "{{ pod_network_cidr }}"
  dnsDomain: "cluster.local"
{% if skip_kube_proxy %}
proxy:
  disabled: true
{% endif %}
apiServer:
  certSANs:
{% for san in api_server_cert_sans %}
    - "{{ san }}"
{% endfor %}
  extraArgs:
    - name: "authorization-mode"
      value: "Node,RBAC"
etcd:
  local:
```

dataDir: "/var/lib/etcd"

apiVersion: kubelet.config.k8s.io/v1beta1

kind: KubeletConfiguration

cgroupDriver: systemd

serverTLSBootstrap: true

4.8 Ansible実行コマンド



bash

cd ansible/

全Playbook実行

ansible-playbook -i inventory/hosts.ini playbooks/01-prerequisites.yml

ansible-playbook -i inventory/hosts.ini playbooks/02-container-runtime.yml

ansible-playbook -i inventory/hosts.ini playbooks/03-kubernetes-install.yml

ansible-playbook -i inventory/hosts.ini playbooks/04-cluster-init.yml

kubeconfigコピー

mkdir -p ~/.kube

cp kubeconfig ~/.kube/config

確認

kubectl get nodes

kubectl get pods -A

5. Cilium CNI設定（kube-proxy完全置換）

5.1 Cilium完全設定

helmfile/values/dev/cilium.yaml:



yaml

Cilium 1.18.2 開発環境設定

クラスタ識別

cluster:
 name: dev-cluster
 id: 1

=====

kube-proxy完全置換

=====

kubeProxyReplacement: true
k8sServiceHost: "192.168.1.100"
k8sServicePort: 6443

eBPF機能

bpf:
 masquerade: true
 lbExternalClusterIP: true

kube-proxy機能

nodePort:
 enabled: true
 mode: hybrid

externalIPs:
 enabled: true

hostPort:
 enabled: true

sessionAffinity: true

=====

LoadBalancer IPAM (MetalLB代替)

=====

l2announcements:
 enabled: true

l2NeighDiscovery:
 enabled: true
 refreshPeriod: "30s"

k8sClientRateLimit:
 qps: 50

burst: 100

=====

Hubble可觀測性

=====

hubble:

enabled: true

relay:

enabled: true

replicas: 1

resources:

requests:

cpu: 50m

memory: 64Mi

limits:

cpu: 500m

memory: 256Mi

ui:

enabled: true

replicas: 1

service:

type: ClusterIP

resources:

requests:

cpu: 50m

memory: 64Mi

limits:

cpu: 500m

memory: 256Mi

metrics:

enabled:

- dns

- drop

- tcp

- flow

- http

serviceMonitor:

enabled: true

port: 9965

=====

Gateway API (HTTP/HTTPS/TLS/gRPC専用)

=====

gatewayAPI:

enabled: true

=====

Prometheus統合

=====

prometheus:

enabled: true

port: 9962

serviceMonitor:

enabled: true

interval: "30s"

Operator

operator:

replicas: 1 # 開発環境は1

prometheus:

enabled: true

serviceMonitor:

enabled: true

Envoy (L7ポリシー)

envoy:

enabled: true

prometheus:

enabled: true

serviceMonitor:

enabled: true

=====

ネットワーク

=====

ipam:

mode: kubernetes

tunnel: vxlan

ipv4:

enabled: true

ipv6:

enabled: false

```
# =====  
# リソース管理（開発環境最適化）  
# =====  
resources:  
  requests:  
    cpu: 100m  
    memory: 256Mi  
  limits:  
    cpu: 2000m  
    memory: 2Gi  
  
# 全ノードデプロイ  
tolerations:  
  - operator: Exists  
  
# ローリングアップデート  
updateStrategy:  
  type: RollingUpdate  
  rollingUpdate:  
    maxUnavailable: 1  
  
rollOutCiliumPods: true  
  
# デバッグ  
debug:  
  enabled: false
```

5.2 LoadBalancer IPプール

cilium-lb-ippool.yaml:



yaml

```
apiVersion: cilium.io/v2alpha1
kind: CiliumLoadBalancerIPPool
metadata:
```

```
  name: dev-pool
```

```
spec:
```

```
  cidrs:
```

```
    - cidr: "192.168.100.0/24"
```

```
---
```

```
apiVersion: cilium.io/v2alpha1
```

```
kind: CiliumL2AnnouncementPolicy
```

```
metadata:
```

```
  name: l2-policy
```

```
  namespace: kube-system
```

```
spec:
```

```
  interfaces:
```

```
    - ^ens[0-9]+
```

```
    - ^eth[0-9]+
```

```
  loadBalancerIPs: true
```

```
  externalIPs: true
```

```
  nodeSelector:
```

```
    matchExpressions:
```

```
      - key: node-role.kubernetes.io/control-plane
```

```
        operator: DoesNotExist
```

適用：



bash

```
kubectl apply -f cilium-lb-ippool.yaml
```

5.3 Cilium確認コマンド



bash

Cilium状態

```
kubectl -n kube-system exec ds/cilium -- cilium status
```

Connectivity Test

```
cilium connectivity test
```

Hubble確認

```
kubectl -n kube-system get pods -l k8s-app=hubble-relay
```

```
kubectl -n kube-system get pods -l k8s-app=hubble-ui
```

LoadBalancer IP確認

```
kubectl get svc -A | grep LoadBalancer
```

6. helmfile環境管理

6.1 helmfile.yaml

helmfile/helmfile.yaml:



yaml

環境定義

environments:

dev:

values:

- environments/dev.yaml

staging:

values:

- environments/staging.yaml

prod:

values:

- environments/prod.yaml

リポジトリ

repositories:

- name: cilium

url: https://helm.cilium.io/

- name: traefik

url: https://traefik.github.io/charts

- name: prometheus-community

url: https://prometheus-community.github.io/helm-charts

- name: grafana

url: https://grafana.github.io/helm-charts

- name: jaegertracing

url: https://jaegertracing.github.io/helm-charts

- name: vector

url: https://helm.vector.dev

- name: minio-operator

url: https://operator.min.io/

- name: argo

url: https://argoproj.github.io/argo-helm

- name: kubernetes-dashboard

url: https://kubernetes.github.io/dashboard/

helmDefaults:

atomic: true

wait: true

timeout: 600

cleanupOnFail: true

recreatePods: false

リリース

releases:

1. Cilium (既にAnsibleでインストール済みの場合はスキップ)

```
# - name: cilium
# namespace: kube-system
# chart: cilium/cilium
# version: 1.18.2
# values:
#   - values/{{ .Environment.Name }}/cilium.yaml
```

2. Traefik

```
- name: traefik
  namespace: traefik
  chart: traefik/traefik
  version: 37.1.2
  createNamespace: true
  values:
    - values/{{ .Environment.Name }}/traefik.yaml
```

3. kube-prometheus-stack

```
- name: kube-prometheus-stack
  namespace: monitoring
  chart: prometheus-community/kube-prometheus-stack
  version: 78.3.0
  createNamespace: true
  values:
    - values/{{ .Environment.Name }}/prometheus.yaml
  needs:
    - traefik/traefik
```

4. Vector

```
- name: vector
  namespace: monitoring
  chart: vector/vector
  version: 0.46.0
  values:
    - values/{{ .Environment.Name }}/vector.yaml
  needs:
    - monitoring/kube-prometheus-stack
```

5. Jaeger Operator

```
- name: jaeger-operator
  namespace: observability
  chart: jaegertracing/jaeger-operator
  version: 2.57.0
  createNamespace: true
  values:
```


- values/{{ .Environment.Name }}/jaeger.yaml

6. MinIO Operator

- name: minio-operator
namespace: minio-operator
chart: minio-operator/operator
version: 7.1.1
createNamespace: true
set:
- name: console.enabled
value: "{{ .Values.minioConsoleEnabled }}"

7. ArgoCD

- name: argocd
namespace: argocd
chart: argo/argo-cd
version: 8.6.2
createNamespace: true
values:
- values/{{ .Environment.Name }}/argocd.yaml
needs:
- traefik/traefik

8. Kubernetes Dashboard

- name: kubernetes-dashboard
namespace: kubernetes-dashboard
chart: kubernetes-dashboard/kubernetes-dashboard
version: 7.13.0
createNamespace: true
values:
- values/{{ .Environment.Name }}/k8s-dashboard.yaml
needs:
- traefik/traefik

6.2 環境別設定

helmfile/environments/dev.yaml:



yaml

開発環境設定

clusterName: dev-cluster

domain: dev.local

controlPlaneEndpoint: "192.168.1.100"

レプリカ数（開発環境最小化）

traefikReplicas: 1

prometheusReplicas: 1

grafanaReplicas: 1

ストレージ

prometheusRetention: 7d

prometheusStorageSize: 20Gi

機能フラグ

minioConsoleEnabled: true

enableServiceMonitor: true

enableIngress: true

認証

enableBasicAuth: false # 開発環境では無効

helmfile/environments/prod.yaml:



yaml

本番環境設定

clusterName: prod-cluster

domain: example.com

controlPlaneEndpoint: "10.0.1.100"

レプリカ数 (HA)

traefikReplicas: 3

prometheusReplicas: 2

grafanaReplicas: 2

ストレージ

prometheusRetention: 30d

prometheusStorageSize: 100Gi

機能フラグ

minioConsoleEnabled: false

enableServiceMonitor: true

enableIngress: true

認証

enableBasicAuth: true

7. 全コンポーネント設定詳細

7.1 Traefik (TCP/UDP対応)

helmfile/values/dev/traefik.yaml:



yaml

Traefik 3.5開発環境設定

デプロイメント

deployment:

replicas: 1

Providers

providers:

kubernetesCRD:

enabled: true

allowCrossNamespace: true

kubernetesIngress:

enabled: true

publishedService:

enabled: true

エントリーポイント

ports:

web:

port: 80

protocol: TCP

expose:

default: true

websecure:

port: 443

protocol: TCP

expose:

default: true

tls:

enabled: true

TCP/UDPポート (CoreDNS, Vector等)

dns-tcp:

port: 5353

protocol: TCP

expose:

default: true

dns-udp:

port: 5353

protocol: UDP

expose:

default: true

vector-tcp:
port: 6000
protocol: TCP
expose:
default: true

Service

service:
type: LoadBalancer # Cilium LB IPAMが割り当て
annotations:
io.cilium/lb-ipam-ips: "192.168.100.10"

リソース

resources:
requests:
cpu: 100m
memory: 128Mi
limits:
cpu: 1000m
memory: 512Mi

メトリクス

metrics:
prometheus:
enabled: true
service:
enabled: true
serviceMonitor:
enabled: true

ログ

logs:
general:
level: INFO
access:
enabled: true
format: json

ダッシュボード (開発環境のみ)

ingressRoute:
dashboard:
enabled: true

```
matchRule: Host(`traefik.dev.local`)
entryPoints: ["web"]
```

7.2 kube-prometheus-stack

helmfile/values/dev/prometheus.yaml:



yaml

kube-prometheus-stack開発環境設定

Prometheus

prometheus:

enabled: true

prometheusSpec:

replicas: 1

retention: 7d

retentionSize: "15GB"

ストレージ

storageSpec:

volumeClaimTemplate:

spec:

accessModes: ["ReadWriteOnce"]

resources:

requests:

storage: 20Gi

ServiceMonitor自動検出

serviceMonitorSelectorNilUsesHelmValues: false

podMonitorSelectorNilUsesHelmValues: false

リソース

resources:

requests:

cpu: 200m

memory: 1Gi

limits:

cpu: 2000m

memory: 4Gi

Grafana

grafana:

enabled: true

replicas: 1

adminPassword: "admin" *# 開発環境*

データソース

additionalDataSources:

- name: Prometheus

type: prometheus

uid: prometheus

access: proxy

url: http://kube-prometheus-stack-prometheus:9090
isDefault: true
jsonData:
 httpMethod: POST
 interval: 30s

- name: Jaeger
type: jaeger
uid: jaeger
access: proxy
url: http://jaeger-query.observability:16686

ダッシュボード

dashboardProviders:
dashboardproviders.yaml:
 apiVersion: 1
 providers:
 - name: 'default'
 orgId: 1
 folder: ''
 type: file
 disableDeletion: false
 editable: true
 options:
 path: /var/lib/grafana/dashboards/default

dashboards:
default:
 node-exporter:
 gnetId: 1860
 revision: 31
 datasource: Prometheus
 kubernetes-cluster:
 gnetId: 7249
 revision: 1
 datasource: Prometheus
 kubernetes-pods:
 gnetId: 6417
 revision: 1
 datasource: Prometheus
 cilium-metrics:
 gnetId: 16611
 revision: 1
 datasource: Prometheus

Ingress

ingress:

enabled: true

ingressClassName: traefik

hosts:

- grafana.dev.local

AlertManager

alertmanager:

enabled: true

node-exporter

nodeExporter:

enabled: true

kube-state-metrics

kubeStateMetrics:

enabled: true

kubelet

kubelet:

enabled: true

serviceMonitor:

interval: 30s

7.3 Vector

helmfile/values/dev/vector.yaml:



yaml

Vector開発環境設定

role: Agent

ConfigMap

customConfig:

data_dir: /var/lib/vector

sources:

Kubernetesログ

kubernetes_logs:

type: kubernetes_logs

ホストメトリクス

host_metrics:

type: host_metrics

filesystem:

devices:

excludes: ["binfmt_misc"]

Vector内部メトリクス

internal_metrics:

type: internal_metrics

transforms:

ラベル追加

add_labels:

type: remap

inputs: ["host_metrics", "internal_metrics"]

source: |

.tags.environment = "dev"

.tags.cluster = "dev-cluster"

sinks:

Prometheus Remote Write

prometheus:

type: prometheus_remote_write

inputs: ["add_labels"]

endpoint: "http://kube-prometheus-stack-prometheus.monitoring:9090/api/v1/write"

compression: snappy

healthcheck:

enabled: true

コンソール出力 (開発)

```
console:
  type: console
  inputs: ["kubernetes_logs"]
  encoding:
    codec: json
```

リソース

```
resources:
  requests:
    cpu: 100m
    memory: 128Mi
  limits:
    cpu: 1000m
    memory: 512Mi
```

7.4 Jaeger

helmfile/values/dev/jaeger.yaml:



yaml

Jaeger Operator開発環境設定

```
jaeger:
  create: true
  spec:
    strategy: allInOne
    allInOne:
      image: jaegertracing/all-in-one:1.65.0
      options:
        memory:
          max-traces: 10000

    storage:
      type: memory

    ingress:
      enabled: true
      ingressClassName: traefik
      hosts:
        - jaeger.dev.local

    agent:
      strategy: DaemonSet
```

7.5 ArgoCD

helmfile/values/dev/argocd.yaml:



yaml

ArgoCD開発環境設定

global:

domain: argocd.dev.local

server:

replicas: 1

Traefik経由アクセス用

extraArgs:

- --insecure # TLS終端はTraefikで行う

ingress:

enabled: true

ingressClassName: traefik

hosts:

- argocd.dev.local

metrics:

enabled: true

serviceMonitor:

enabled: true

repoServer:

replicas: 1

controller:

replicas: 1

metrics:

enabled: true

serviceMonitor:

enabled: true

redis:

enabled: true

Dex (SSO) は開発環境で無効

dex:

enabled: false

初期パスワード

configs:

secret:

argocdServerAdminPassword: "\$2a\$10\$..." # bcrypt hash

7.6 Kubernetes Dashboard

helmfile/values/dev/k8s-dashboard.yaml:



yaml

Kubernetes Dashboard開発環境設定

app:

ingress:

enabled: true

ingressClassName: traefik

hosts:

- k8s-dashboard.dev.local

HTTPSバックエンド用

annotations:

traefik.ingress.kubernetes.io/router.tls: "true"

メトリクススクレイパー

metricsScraper:

enabled: true

リソース

resources:

requests:

cpu: 100m

memory: 128Mi

limits:

cpu: 500m

memory: 512Mi

7.7 MinIO

minio-tenant.yaml（手動作成）:



yaml

```
apiVersion: v1
kind: Namespace
metadata:
  name: minio
---
apiVersion: v1
kind: Secret
metadata:
  name: storage-config
  namespace: minio
stringData:
  config.env: |
    export MINIO_ROOT_USER="admin"
    export MINIO_ROOT_PASSWORD="DevPassword123"
---
apiVersion: minio.min.io/v2
kind: Tenant
metadata:
  name: dev-storage
  namespace: minio
spec:
  image: quay.io/minio/minio:latest
  pools:
    - name: pool-0
      servers: 4
      volumesPerServer: 2
      volumeClaimTemplate:
        spec:
          accessModes:
            - ReadWriteOnce
          resources:
            requests:
              storage: 50Gi
  configuration:
    name: storage-config
  requestAutoCert: false
  buckets:
    - name: prometheus
    - name: loki
    - name: grafana
```

8. メトリクス可視化設定

8.1 主要メトリクス一覧

ノードメトリクス (node-exporter)

- node_cpu_seconds_total - CPU使用時間
- node_memory_MemAvailable_bytes - 利用可能メモリ
- node_memory_MemTotal_bytes - 総メモリ
- node_filesystem_avail_bytes - ディスク空き容量
- node_network_receive_bytes_total - ネットワーク受信
- node_network_transmit_bytes_total - ネットワーク送信

Podメトリクス (cAdvisor)

- container_cpu_usage_seconds_total - コンテナCPU
- container_memory_working_set_bytes - コンテナメモリ
- container_network_receive_bytes_total - コンテナRX
- container_network_transmit_bytes_total - コンテナTX

Ciliumメトリクス

- cilium_datapath_conntrack_gc_runs_total - Connection tracking
- cilium_drop_count_total - パケットドロップ
- cilium_forward_count_total - パケット転送
- cilium_policy_l7_total - L7ポリシー適用

8.2 Traefik IngressRoute設定

ingress-routes.yaml:



yaml

Grafana

apiVersion: traefik.io/v1alpha1

kind: IngressRoute

metadata:

name: grafana

namespace: monitoring

spec:

entryPoints:

- web

routes:

- kind: Rule

match: Host(`grafana.dev.local`)

services:

- name: kube-prometheus-stack-grafana

port: 80

Prometheus

apiVersion: traefik.io/v1alpha1

kind: IngressRoute

metadata:

name: prometheus

namespace: monitoring

spec:

entryPoints:

- web

routes:

- kind: Rule

match: Host(`prometheus.dev.local`)

services:

- name: kube-prometheus-stack-prometheus

port: 9090

Hubble UI

apiVersion: traefik.io/v1alpha1

kind: IngressRoute

metadata:

name: hubble-ui

namespace: kube-system

spec:

entryPoints:

- web

routes:

- kind: Rule

match: Host(`hubble.dev.local`)

services:

- name: hubble-ui

port: 80

ArgoCD UI

apiVersion: traefik.io/v1alpha1

kind: IngressRoute

metadata:

name: argocd-server

namespace: argocd

spec:

entryPoints:

- web

routes:

- kind: Rule

match: Host(`argocd.dev.local`)

services:

- name: argocd-server

port: 80

Jaeger UI

apiVersion: traefik.io/v1alpha1

kind: IngressRoute

metadata:

name: jaeger-query

namespace: observability

spec:

entryPoints:

- web

routes:

- kind: Rule

match: Host(`jaeger.dev.local`)

services:

- name: jaeger-query

port: 16686

Kubernetes Dashboard

apiVersion: traefik.io/v1alpha1

```
kind: IngressRoute
metadata:
  name: k8s-dashboard
  namespace: kubernetes-dashboard
spec:
  entryPoints:
    - web
  routes:
    - kind: Rule
      match: Host(`k8s-dashboard.dev.local`)
      services:
        - name: kubernetes-dashboard-kong-proxy
          port: 443
          scheme: https
          serversTransport: dashboard-transport
```

```
---
# HTTPS Backend用トランスポート
apiVersion: traefik.io/v1alpha1
kind: ServersTransport
metadata:
  name: dashboard-transport
  namespace: kubernetes-dashboard
spec:
  serverName: kubernetes-dashboard
  insecureSkipVerify: true
```

8.3 hosts設定（ローカル開発）



bash

```
# /etc/hostsに追加
192.168.100.10 grafana.dev.local
192.168.100.10 prometheus.dev.local
192.168.100.10 hubble.dev.local
192.168.100.10 argocd.dev.local
192.168.100.10 jaeger.dev.local
192.168.100.10 k8s-dashboard.dev.local
192.168.100.10 traefik.dev.local
```

9. デプロイ手順

9.1 完全デプロイフロー



bash

```
# =====
```

```
# Step 1: Terraformでインフラ構築
```

```
# =====
```

```
cd terraform/
```

```
terraform init
```

```
terraform apply -auto-approve
```

```
# =====
```

```
# Step 2: Ansibleでクラスタ構築
```

```
# =====
```

```
cd ../ansible/
```

```
# 全Playbook実行
```

```
ansible-playbook -i inventory/hosts.ini playbooks/01-prerequisites.yml
```

```
ansible-playbook -i inventory/hosts.ini playbooks/02-container-runtime.yml
```

```
ansible-playbook -i inventory/hosts.ini playbooks/03-kubernetes-install.yml
```

```
ansible-playbook -i inventory/hosts.ini playbooks/04-cluster-init.yml
```

```
# kubeconfigコピー
```

```
mkdir -p ~/.kube
```

```
cp kubeconfig ~/.kube/config
```

```
chmod 600 ~/.kube/config
```

```
# クラスタ確認
```

```
kubectl get nodes
```

```
kubectl get pods -A
```

```
# =====
```

```
# Step 3: Cilium LB IPAMセットアップ
```

```
# =====
```

```
kubectl apply -f cilium-lb-ippool.yaml
```

```
# Cilium確認
```

```
kubectl -n kube-system exec ds/cilium -- cilium status
```

```
# =====
```

```
# Step 4: helmfileで全コンポーネントデプロイ
```

```
# =====
```

```
cd ../helmfile/
```

```
# 開発環境デプロイ
```

```
helmfile -e dev sync
```

```
# 進捗確認
```

`watch kubectl get pods -A`

`# =====`

`# Step 5: MinIO Tenantデプロイ`

`# =====`

`kubectl apply -f minio-tenant.yaml`

`# =====`

`# Step 6: IngressRoute設定`

`# =====`

`kubectl apply -f ingress-routes.yaml`

`# =====`

`# Step 7: 確認`

`# =====`

`# 全Pod確認`

`kubectl get pods -A`

`# Service確認 (LoadBalancer IP)`

`kubectl get svc -A | grep LoadBalancer`

`# IngressRoute確認`

`kubectl get ingressroute -A`

`# メトリクス確認`

`kubectl top nodes`

`kubectl top pods -A`

`# UIアクセス`

`echo "Grafana: http://grafana.dev.local"`

`echo "Prometheus: http://prometheus.dev.local"`

`echo "Hubble: http://hubble.dev.local"`

`echo "ArgoCD: http://argocd.dev.local"`

`echo "Jaeger: http://jaeger.dev.local"`

`echo "K8s Dashboard: http://k8s-dashboard.dev.local"`

9.2 デプロイ時間目安

フェーズ	所要時間
Terraform VM作成	5-10分
Ansible前提条件	5分
Container Runtime	3分
Kubernetesインストール	5分
クラスタ初期化 + Cilium	10分
helmfile全コンポーネント	15-20分
合計	40-50分

10. トラブルシューティング

10.1 Cilium関連

問題: ノードがNotReady



bash

```
# Cilium Pod確認
kubectl get pods -n kube-system -l k8s-app=cilium

# ログ確認
kubectl logs -n kube-system ds/cilium --tail=50

# Cilium状態
kubectl -n kube-system exec ds/cilium -- cilium status

# Connectivity Test
cilium connectivity test
```

問題: LoadBalancer IPが割り当たらない



bash

IPプール確認

```
kubectl get ciliumloadbalancerippool
```

L2アナウンス確認

```
kubectl get ciliuml2announcementpolicy
```

Ciliumログでエラー確認

```
kubectl logs -n kube-system deploy/cilium-operator | grep -i "lb-ipam"
```

10.2 Prometheus関連

問題: メトリクスが表示されない



bash

ServiceMonitor確認

```
kubectl get servicemonitor -A
```

Prometheusターゲット確認 (Port Forward)

```
kubectl port-forward -n monitoring svc/kube-prometheus-stack-prometheus 9090:9090
```

```
# http://localhost:9090/targets
```

Prometheusログ

```
kubectl logs -n monitoring sts/prometheus-kube-prometheus-stack-prometheus
```

問題: Grafanaダッシュボードが空



bash

データソース確認

```
kubectl get cm -n monitoring kube-prometheus-stack-grafana -o yaml | grep -A 10 datasources
```

Grafanaログ

```
kubectl logs -n monitoring deploy/kube-prometheus-stack-grafana
```

10.3 Traefik関連

問題: IngressRouteが動作しない



bash

```
# IngressRoute確認
kubectl get ingressroute -A

# Traefikログ
kubectl logs -n traefik deploy/traefik

# Traefik設定ダンプ
kubectl port-forward -n traefik svc/traefik 9000:9000
# http://localhost:9000/dashboard/
```

10.4 ArgoCD関連

問題: 502 Bad Gateway



bash

```
# ArgoCD server.insecure設定確認
kubectl get cm argocd-cmd-params-cm -n argocd -o yaml | grep insecure

# 設定追加
kubectl patch cm argocd-cmd-params-cm -n argocd \
--type merge \
-p '{"data":{"server.insecure":"true"}}'

kubectl rollout restart deployment argocd-server -n argocd
```

10.5 一般的な問題

問題: Podが起動しない



bash

Pod詳細

```
kubectl describe pod <pod-name> -n <namespace>
```

イベント確認

```
kubectl get events -n <namespace> --sort-by='.lastTimestamp'
```

ログ確認

```
kubectl logs <pod-name> -n <namespace> --previous
```

問題: ノードリソース不足



bash

リソース使用状況

```
kubectl top nodes
```

```
kubectl describe node <node-name>
```

Pod退避

```
kubectl drain <node-name> --ignore-daemonsets --delete-emptydir-data
```

10.6 完全リセット手順



bash

クラスタリセット

```
ansible-playbook -i inventory/hosts.ini playbooks/reset-cluster.yml
```

または手動

```
sudo kubeadm reset -f
```

```
sudo rm -rf /etc/kubernetes /var/lib/etcd /var/lib/kubelet /var/lib/cni
```

```
sudo iptables -F && sudo iptables -t nat -F && sudo iptables -t mangle -F && sudo iptables -X
```

11. 次のステップ

11.1 本番環境への移行

1. Control Plane 3台構成へ拡張（etcd quorum確保）
2. 外部ロードバランサー追加（HAProxy/keepalived）

3. **cert-manager**導入（Let's Encrypt自動化）
4. **Velero**バックアップ戦略
5. **OPA Gatekeeper**ポリシー管理

11.2 GitOps完全移行



bash

```
# ArgoCD Application作成
kubect apply -f - <<EOF
apiVersion: argoproj.io/v1alpha1
kind: Application
metadata:
  name: infrastructure
  namespace: argocd
spec:
  project: default
  source:
    repoURL: https://github.com/yourorg/k8s-infrastructure
    targetRevision: main
    path: helmfile
  destination:
    server: https://kubernetes.default.svc
    namespace: default
  syncPolicy:
    automated:
      prune: true
      selfHeal: true
EOF
```

11.3 監視強化

- **Thanos**導入（Prometheus長期保存）
- **Loki**導入（ログ集約）
- **Tempo**導入（トレース保存）
- **アラートルール**カスタマイズ

まとめ

本ガイドでは、Proxmox + Terraform + Ansible + helmfileを用いた**Kubernetes 1.33基盤の完全自動構築**を解説しました。

重要なポイント：

- 1. ☒ **Cilium kube-proxy完全置換**でeBPF高性能化
- 2. ☒ **Cilium LB IPAM**でMetalLB不要
- 3. ☒ **Cilium Gateway API**はHTTP/HTTPS/TLS/gRPC専用（TCP/UDPは未実装）
- 4. ☒ **Traefik**でTCP/UDPルーティング対応
- 5. ☒ **helmfile環境管理**で開発/本番環境切り替え
- 6. ☒ **ノード・Podメトリクス完全可視化**（Prometheus + Grafana）
- 7. ☒ **全UI統一アクセス**（Traefik IngressRoute）

開発環境構成：Control Plane 2台 + Worker 3台でリソース最適化済み。本番環境へはhelm valuesとenvironments YAMLの変更のみで移行可能です。

デプロイ時間：約40-50分で完全自動構築が完了します。

作成日: 2025年10月19日

対象バージョン: Kubernetes 1.33.5, Cilium 1.18.2

次回更新: Kubernetes 1.34正式対応、Cilium TCP/UDPRoute実装時